

MATH 250: Mathematical Data Visualization

Course overview: data as representation

Peter A. Gao

January 24, 2024

San José State University

What is the focus of this course?

Key questions

- **How do we use vectors and matrices to represent images/networks/other forms of information as data?**
 - *data as representation*
- **How can we summarize and visualize high-dimensional datasets?**
 - *linear algebra, dimensionality reduction, data visualization*
- **How do we extract insights from large, often unstructured, datasets?**
 - *pattern recognition, data mining, classification, clustering*

Example: handwritten digits




Figure 1: Handwritten digits from MNIST dataset [1]

Example: handwritten digits

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]	[,14]	[,15]	[,16]	[,17]	[,18]	[,19]	[,20]	[,21]	[,22]	[,23]	[,24]	[,25]	[,26]	[,27]	[,28]
[1,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[2,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[3,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[4,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[5,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[6,]	0	0	0	0	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	247	127	0	0	0	0
[7,]	0	0	0	0	0	0	0	0	30	36	94	154	170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0
[8,]	0	0	0	0	0	0	0	49	238	253	253	253	253	253	253	253	251	93	82	82	56	39	0	0	0	0	0	0
[9,]	0	0	0	0	0	0	0	18	219	253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	0
[10,]	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154	0	0	0	0	0	0	0	0	0	0	0
[11,]	0	0	0	0	0	0	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[12,]	0	0	0	0	0	0	0	0	0	139	253	190	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[13,]	0	0	0	0	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[14,]	0	0	0	0	0	0	0	0	0	0	35	241	225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0
[15,]	0	0	0	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	0	0	0	0	0	0	0	0
[16,]	0	0	0	0	0	0	0	0	0	0	0	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0
[17,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	0	0	0	0	0	0	0	0	0
[18,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	0
[19,]	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	253	207	2	0	0	0	0	0	0	0	0
[20,]	0	0	0	0	0	0	0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0	0	0	0	0
[21,]	0	0	0	0	0	0	0	0	0	24	114	221	253	253	253	201	78	0	0	0	0	0	0	0	0	0	0	0
[22,]	0	0	0	0	0	0	0	23	66	213	253	253	253	198	81	2	0	0	0	0	0	0	0	0	0	0	0	0
[23,]	0	0	0	0	0	0	18	171	219	253	253	253	195	80	9	0	0	0	0	0	0	0	0	0	0	0	0	0
[24,]	0	0	0	0	55	172	226	253	253	253	253	244	133	11	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[25,]	0	0	0	0	136	253	253	253	212	135	132	16	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[26,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[27,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[28,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

Figure 2: Handwritten digit from MNIST dataset, as matrix

Example: handwritten digits

Each handwritten digit is digitized and represented using a 28 x 28 pixel grid of numbers between 0 (white) and 255 (black).

We can represent each image using a 28 x 28 matrix, or a 1 x 784 row vector. In other words, each image is a point in the 784-dimensional space \mathbb{R}^{784} .

This course focuses on the task of **dimensionality reduction**. Many modern datasets have hundreds/thousands/millions of dimensions/features/variables.

Reducing dimension **can**:

- save run time and storage space
- reduce noise and multi-collinearity, facilitating statistical analysis
- facilitate data visualization

Example: handwritten digits

What if we wish to develop an algorithm for classifying our digits?

Working in the original 784-dimensional space may not be necessary (for example, the corners of each image are almost always 0 and not very informative).

Feature selection approaches examine a subset of the original dimensions/features.

- *ex. subset selection, ridge regression, LASSO*

This course will focus on **feature extraction**: constructing a low-dimensional representation of the original high-dimensional data

- *ex. principal component analysis, manifold learning*

Example: handwritten digits

Instead of working in \mathbb{R}^{784} , what if we work in a lower-dimensional space?

In this class, we'll learn what it means to construct a lower-dimensional representation of a high-dimensional object and then visualize our representations.

Aside: data as representation

Note that when we digitize handwritten digits, we must choose a resolution (28 x 28 in this case). The higher the resolution, the more information is contained in our digital representation.

In this sense, our data are typically just representations of some real world objects. Lowering the "resolution" of our representation loses information but typically saves space and can simplify analysis.

Aside: The "right" resolution



Figure 3: Coastline paradox, illustrated by [Alexandre Van de Sande](#).

Dimensionality reduction facilitates many common statistical and machine learning tasks:

- **visualization**
- regression (MATH 261A)
- classification (MATH 251)
- clustering (MATH 252)

We will touch on all of these, but our focus will be on linear algebra and visualization tools that are foundational for machine learning.

This course covers **matrix methods and linear algebra** for machine learning and high-dimensional data analysis.

Our focus will be on **dimension reduction**, as applied to **data visualization**.

We will also cover computation and tools for visualization in R.

This course is used in the following ways:

- Prerequisite for **MATH 251 Statistical and Machine Learning Classification**
- Elective for the regular **MS Statistics** program
- Required for the **MS Statistics Machine Learning Specialization**
- Required for the **MS Data Science** program (joint with CS)

Prerequisites:

- **Math 32:** Multivariate calculus
- **Math 39:** Linear algebra
- **Math 163:** Probability theory

In addition, though there is no formal programming prerequisite, you are expected to solve problems and make visualizations using R.

Coursework

- **Homework (40%):** Roughly two weeks per assignment, involving both problem sets (proofs) and R labs.
- **Exams:** Two midterm exams (**20% each**)
- **Final project (20%):** Final report and presentation during last week of class

In general, the late policy is as follows:

- Any assignment that is received late but less than 24 hours late will receive a penalty of 25%.
- Any assignment that is received 24–48 hours late will receive a penalty of 50%.
- Assignments will not be accepted more than 48 hours late.

Extensions for academic purposes (ex. conference presentation or job interview) or extreme circumstances (ex. illness, emergency) will generally be granted. Email me at least 24 hours before an assignment is due to request an extension, along with the reason.

- You may discuss problems, approaches, and solutions with your classmates.
- You must credit anyone with whom you worked on each assignment.
- All submitted work must be your own; you should not submit code or answers copied from any resource including your classmates.
- If you have any questions, ask!

External resources

Students may use external online resources including discussion forums (ex. StackOverflow) large language model-based chatbots (ex. ChatGPT) as aids for learning and understanding course material.

However, submitting code or answers for course assignments obtained using resources like StackOverflow or ChatGPT is **not permitted**.

When external resources are consulted for an assignment, they must be cited clearly at the top of your submission.

There will be two in-class midterm exams, both closed note and closed book.

The second midterm will be cumulative (essentially an early final).

These exams will test your understanding of the **linear algebra theory and technical details of the dimension reduction methods** covered in class.

Practice problems will be provided before the exams.

Final project

This course culminates with a project in which you will explore a dataset or dimension reduction technique of your choice.

You will prepare a brief oral presentation and write a final report.

Examples:

- Present a dimension reduction method not covered in class.
- Apply a dimension reduction method covered in this course to a real-world dataset and interpret the results.
- Compare two dimension reduction methods using a real-world dataset or simulated data.
- Extend an existing dimension reduction method or review a theoretical result/paper not discussed in class.

Where to look

If you ever have questions about materials, please contact me.
If you need any kind of accommodations, please let me know as soon as possible.

- **Course website:** Course slides, assignment instructions.
- **Canvas:** Official syllabus, receiving grades, data.
- **Gradescope:** Submitting assignments.

Required Texts:

- Strang, G. (2019). Linear Algebra and Learning From Data. Available for free via [Library Course Reserves](#).
- Murphy, K. P. (2022). Probabilistic Machine Learning: An Introduction. Available for free via [Library Course Reserves](#).
- Murphy, K. P. (2023). Probabilistic Machine Learning: Advanced Topics. Available for free via [Library Course Reserves](#).

Suggested texts:

- Baker, J. D. (2024). Applied Multivariate Statistics in R. Available for free [here](#).
- Healy, K. (2019). Data Visualization: A Practical Introduction. Draft version available [here](#).
- Strang, G. (2023). Introduction to Linear Algebra, Sixth Edition.

Texts:

- **Math 32:** Multivariate calculus
- **Math 39:** Linear algebra
- **Math 163:** Probability theory

In addition, though there is no formal programming prerequisite, you are expected to solve problems and make visualizations using R.

Your responsibilities

This course may be challenging and demanding. We are combining difficult theory (linear algebra) and applications (dimension reduction and visualization).

However, this material is crucial for establishing the theoretical and applied foundations of machine learning and high-dimensional statistics.

As such, you should strive to:

- Attend all classes
- Participate in class discussion
- Read the textbooks before and after class
- Think carefully through homework problems
- **Ask questions early and often**

Office hours: MW 10:30-11:30 or by appointment in MH311 (or Zoom, by appointment)

Canvas: Discussion board; up to 2% extra credit for participation on the discussion board.

Email: peter.gao [at sjsu]. Feel free to send me a reminder after 48 hours have passed. Please include [MATH 250] in your subject line.

Getting to know each other

Form groups of 4-5 and discuss the following:

- Introduce yourself (names, major/program)
- What are you excited/nervous/confused about with regards to this course? What questions do you have?
- Have you ever used R? Programmed?
- Why did you choose to study statistics/data science? What is one area of interest you would like to use statistics/data science to study?



Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner,
“Gradient-based learning applied to document recognition,”
Proceedings of the IEEE, vol. 86, pp. 2278–2324, Nov. 1998.