

MATH 250: Mathematical Data Visualization

Applications: Classification

Peter A. Gao

2024-04-15

San José State University

- Classification via linear discriminant analysis (based on lecture by [Guangliang Chen](#))

- Chapter on LDA (applied)
- PSU Stat 508

Linear discriminant analysis

Recall that PCA is a linear dimension reduction technique that identifies the subspace spanned by the maximum-variance directions.


When performing classification, we may want to identify a lower dimensional representation that preserves differences between groups.



Linear discriminant analysis

Linear discriminant analysis seeks to identify the direction that best **separates** the various classes of the dataset.

Note that projections onto parallel lines yield the same separation.



Two-class linear discriminant analysis

Objective: Given data $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^p$ belonging to disjoint classes C_1 and C_2 , identify a line that best separates the classes:

$$\mathbf{w}(t) = t\mathbf{w} + \mathbf{b}, \quad t \in \mathbb{R}$$

where $\mathbf{w}, \mathbf{b} \in \mathbb{R}^p$ and $\|\mathbf{w}\| = 1$.

Two-class linear discriminant analysis

Since parallel lines have the same separation, we focus on lines passing through the origin:

$$\mathbf{w}(t) = t\mathbf{w}, \quad t \in \mathbb{R}$$


where $\|\mathbf{w}\| = 1$. Consider that the projections of the data onto this line are given by

$$\mathbf{p}_i = (\mathbf{x}_i^\top \mathbf{w})\mathbf{w} := a_i\mathbf{w}, \quad i = 1, \dots, n$$

Two-class linear discriminant analysis

Given a projection onto a line, how can we quantify the “amount of separation” between the classes?

Our goal will be to identify the direction that yields the best separation.



Two-class linear discriminant analysis

One possibility is simply to measure the distances between the two means along the line: $|\mu_1 - \mu_2|$, where

$$\mu_1 = \frac{1}{n_1} \sum_{i \in C_1} \mathbf{w}^\top \mathbf{x}_i = \mathbf{w}^\top \left(\frac{1}{n_1} \sum_{i \in C_1} \mathbf{x}_i \right) = \mathbf{w}^\top \mathbf{m}_1$$

where

$$\mathbf{m}_1 = \frac{1}{n_1} \sum_{i \in C_1} \mathbf{x}_i$$

Similarly,

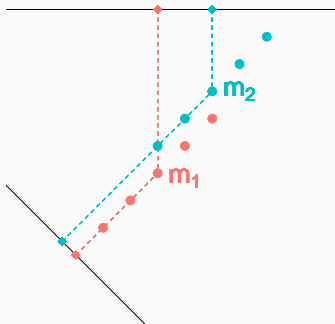
$$\mu_2 = \mathbf{w}^\top \mathbf{m}_2, \quad \mathbf{m}_2 = \frac{1}{n_2} \sum_{i \in C_2} \mathbf{x}_i$$

Two-class linear discriminant analysis

In other words, the problem becomes

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} |\mu_1 - \mu_2| = \max_{\mathbf{w}: \|\mathbf{w}\|=1} |\mathbf{w}^\top (\mathbf{m}_1 - \mathbf{m}_2)|$$

However, it turns out this problem does not necessarily work well, since separating the means does not always separate the data well.



Two-class linear discriminant analysis

One way to improve this criterion is to consider the **within-class variances** of the projections:

$$s_1^2 = \sum_{i \in C_1} (\mathbf{w}^\top \mathbf{x}_i - \mu_1)^2, \quad s_2^2 = \sum_{i \in C_2} (\mathbf{w}^\top \mathbf{x}_i - \mu_i)^2$$

Ideally the direction used will ensure that the class means are well-separated (maximizing $(\mu_1 - \mu_2)^2$), while also minimizing the within-class variances s_1^2 and s_2^2 :

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{(\mu_1 - \mu_2)^2}{s_1^2 + s_2^2}$$

Obtaining a solution

How do we solve this optimization problem? First, consider that

$$\begin{aligned}(\mu_1 - \mu_2)^2 &= (\mathbf{w}^\top \mathbf{m}_1 - \mathbf{w}^\top \mathbf{m}_2)^2 \\&= (\mathbf{w}^\top (\mathbf{m}_1 - \mathbf{m}_2))^2 \\&= \mathbf{w}^\top (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^\top \mathbf{w} \\&:= \mathbf{w}^\top S_b \mathbf{w}\end{aligned}$$

where

$$S_b = (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^\top \in \mathbb{R}^{p \times p}$$

is called the between-class scatter matrix. What can we say about S_b ?

Observe that S_b is square, symmetric, and positive semidefinite.

In addition, S_b is rank one, which in turn implies that it has only one positive eigenvalue.

Obtaining a solution

For each class j , the variance of the projections is

$$\begin{aligned}s_j^2 &= \sum_{i \in C_j} (\mathbf{w}^\top \mathbf{x}_i - \mu_j)^2 \\&= \sum_{i \in C_j} (\mathbf{w}^\top \mathbf{x}_i - \mathbf{w}^\top \mathbf{m}_j)^2 \\&= \sum_{i \in C_j} \mathbf{w}^\top (\mathbf{x}_i - \mathbf{m}_j) (\mathbf{x}_i - \mathbf{m}_j)^\top \mathbf{w} \\&= \mathbf{w}^\top \left(\sum_{i \in C_j} (\mathbf{x}_i - \mathbf{m}_j) (\mathbf{x}_i - \mathbf{m}_j)^\top \right) \mathbf{w} \\&= \mathbf{w}^\top S_j \mathbf{w}\end{aligned}$$

where

$$S_j = \sum_{i \in C_j} (\mathbf{x}_i - \mathbf{m}_j) (\mathbf{x}_i - \mathbf{m}_j)^\top$$

is called the within-class scatter matrix for class j .

Obtaining a solution

Thus,

$$s_1^2 + s_2^2 = \mathbf{w}^\top S_1 \mathbf{w} + \mathbf{w}^\top S_2 \mathbf{w} = \mathbf{w}^\top (S_1 + S - 2) \mathbf{w} = \mathbf{w}^\top S_w \mathbf{w}$$

where

$$S_w = S_1 + S_2 = \sum_j \sum_{i \in C_j} (\mathbf{x}_i - \mathbf{m}_j)(\mathbf{x}_i - \mathbf{m}_j)^\top$$

is called the **total within-class scatter matrix** of the data.

What do we know about S_w ?

Obtaining a solution

Thus,

$$s_1^2 + s_2^2 = \mathbf{w}^\top S_1 \mathbf{w} + \mathbf{w}^\top S_2 \mathbf{w} = \mathbf{w}^\top (S_1 + S - 2) \mathbf{w} = \mathbf{w}^\top S_w \mathbf{w}$$

where

$$S_w = S_1 + S_2 = \sum_j \sum_{i \in C_j} (\mathbf{x}_i - \mathbf{m}_j)(\mathbf{x}_i - \mathbf{m}_j)^\top$$

is called the **total within-class scatter matrix** of the data.

What do we know about S_w ?

S_w is square, symmetric, and positive semidefinite.

Thus,

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{(\mu_1 - \mu_2)^2}{s_1^2 + s_2^2} = \max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{\mathbf{w}^\top S_b \mathbf{w}}{\mathbf{w}^\top S_w \mathbf{w}}$$

Look familiar? When S_w is nonsingular, this is a **generalized Rayleigh quotient**.

Theorem

If S_w is nonsingular, the above optimization problem is solved by the generalized eigenvector \mathbf{w}_1 of (S_b, S_w) :

$$S_b \mathbf{w} = \lambda_1 S_w \mathbf{w}_1 \iff S_w^{-1} S_b \mathbf{w}_1 = \lambda_1 \mathbf{w}_1$$

where λ_1 is the largest generalized eigenvalue of (S_b, S_w) .

Review: Generalized Rayleigh quotients

Proposed solution:

$$S_b \mathbf{w}_1 = \lambda_1 S_w \mathbf{w}_1 \iff S_w^{-1} S_b \mathbf{w}_1 = \lambda_1 \mathbf{w}_1$$

Proof: Since S_b is positive definite, it has a square root $S_w^{1/2}$. If $\mathbf{y} = S_w^{1/2} \mathbf{w}$, then

$$\mathbf{w}^\top S_w \mathbf{w} = \mathbf{y}^\top \mathbf{y}$$

Moreover,

$$\mathbf{w}^\top S_b \mathbf{w} = \mathbf{y}^\top (S_w^{-1/2})^\top S_b S_w^{-1/2} \mathbf{y}$$

so the generalized Rayleigh quotient problem is now

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{\mathbf{w}^\top S_b \mathbf{w}}{\mathbf{w}^\top S_w \mathbf{w}} = \max_{\mathbf{y}: \|\mathbf{y}\|=1} \frac{\mathbf{y}^\top (S_w^{-1/2})^\top S_b S_w^{-1/2} \mathbf{y}}{\mathbf{y}^\top \mathbf{y}}$$

Review: Generalized Rayleigh quotients

Thus, we obtain the following solution:

$$\begin{aligned}(S_w^{-1/2})^\top S_b S_w^{-1/2} \mathbf{y} = \lambda_1 \mathbf{y} &\iff (S_w^{-1/2})^\top S_b S_w^{-1/2} S_w^{1/2} \mathbf{w} = \lambda_1 S_w^{1/2} \mathbf{w} \\ &\iff S_w^{-1} S_b \mathbf{w}_1 = \lambda_1 \mathbf{w}_1\end{aligned}$$

Note that for the LDA problem, $\text{rank}(S_w^{-1} S_b) = \text{rank}(S_b) = 1$ so λ_1 is the only non-zero eigenvalue (and it is positive).

Review: Generalized Rayleigh quotients

Alternatively,

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{\mathbf{w}^\top S_b \mathbf{w}}{\mathbf{w}^\top S_w \mathbf{w}}$$

is equivalent to the constrained optimization problem

$$\max_{\mathbf{w}} \mathbf{w}^\top S_b \mathbf{w} \quad \text{subject to} \quad \mathbf{w}^\top S_w \mathbf{w} = 1$$

We can utilize the method of Lagrange multipliers to compute a solution (exercise).

Mathematically, we can solve this problem as follows:

1. Invert the $p \times p$ matrix S_w .
2. Compute $S_w^{-1} S_b$
3. Solve the eigenvalue problem $S_w^{-1} S_b \mathbf{w}_1 = \lambda_1 \mathbf{w}_1$.

This will be computationally expensive (especially inverting S_w).

Alternatively, observe that

$$\begin{aligned}\lambda_1 \mathbf{w}_1 &= S_w^{-1} S_b \mathbf{w}_1 \\ &= S_w^{-1} (\mathbf{m}_1 - \mathbf{m}_2) (\mathbf{m}_1 - \mathbf{m}_2)^\top \mathbf{w}_1\end{aligned}$$

Moreover, $(\mathbf{m}_1 - \mathbf{m}_2)^\top \mathbf{w}_1$ is a scalar. Therefore,

$$\mathbf{w}_1 \propto S_w^{-1} (\mathbf{m}_1 - \mathbf{m}_2)$$

Instead of computing this directly (and inverting S_w), we can solve the linear system

$$S_w^{-1} \mathbf{b} = \mathbf{m}_1 - \mathbf{m}_2$$

for $\mathbf{b} \in \mathbb{R}^p$.

Two-class linear discriminant analysis: summary

The linear discriminant is in the direction

$$\mathbf{w} \propto S_w^{-1}(\mathbf{m}_1 - \mathbf{m}_2)$$

which solves the optimization problem

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{(\mu_1 - \mu_2)^2}{s_1^2 + s_2^2} = \max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{\mathbf{w}^\top S_b \mathbf{w}}{\mathbf{w}^\top S_w \mathbf{w}}$$

where

$$S_b = (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^\top$$
$$S_w = S_1 + S_2 = \sum_j \sum_{i \in C_j} (\mathbf{x}_i - \mathbf{m}_j)(\mathbf{x}_i - \mathbf{m}_j)^\top$$

Dealing with singularity of S_w

Note that solving the generalized eigenvalue problem requires that the **total within-class scatter matrix** S_w is non singular, which yields a solution to

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{(\mu_1 - \mu_2)^2}{s_1^2 + s_2^2} = \max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{\mathbf{w}^\top S_b \mathbf{w}}{\mathbf{w}^\top S_w \mathbf{w}}$$

However, S_w may be singular (or nearly singular), especially if p is large.

This results when the n rows of the data matrix do not span \mathbb{R}^p .

A simple solution is to apply PCA to the data matrix X , yielding principal components Y .

Note that in this case, we do not keep all of the principal components. One rule of thumb is to keep as many PCs as needed to capture 95% of the variation in the data.

We can then perform LDA using the rank- k reduced data matrix.

Another approach is to regularize S_w to create a nonsingular matrix:

$$\begin{aligned} S_w^\beta &= S_w + \beta I_p \\ &= Q\Lambda Q^\top + \beta I_p \\ &= Q(\Lambda + \beta I_p)Q^\top \end{aligned}$$

where $\beta > 0$ is a tuning parameter.

Classification

We can classify new observations by:

1. Projecting the new observations onto the linear discriminant direction(s) $\mathbf{y}_i = \mathbf{w}^\top \mathbf{x}_i$
2. Assume that for each class j , the projected observations $\mathbf{y}_i \sim N(\mu_j, \Sigma_j)$ for $i \in C_j$. This is called the **class conditional distribution**
3. Compute the class posterior

$$p(i \in C_j \mid X, \theta) \propto \pi_j f(\mathbf{y}_i \mid \mu_j, \Sigma_j)$$

where X is the data matrix, θ denotes model parameters, π_c is the **prior probability** of belonging to class j , and $f(\mathbf{y}_i \mid \mu_j, \Sigma_j)$ is the **class conditional density**.

Two-class linear discriminant analysis: example


Consider the following example:

Class 1 has points

$$\{(1, 2), (2, 3), (3, 4.9)\}$$

and Class 2 has points

$$\{(2, 1), (3, 2), (4, 3.9)\}$$



Two-class linear discriminant analysis: example

Then $\mathbf{m}_1 = (2, 3.3)^\top$ and

$$\begin{aligned} S_1 &= \left(\begin{bmatrix} 1 \\ 2 \end{bmatrix} - \begin{bmatrix} 2 \\ 3.3 \end{bmatrix} \right) \left(\begin{bmatrix} 1 & 2 \end{bmatrix} - \begin{bmatrix} 2 & 3.3 \end{bmatrix} \right) \\ &+ \left(\begin{bmatrix} 2 \\ 3 \end{bmatrix} - \begin{bmatrix} 2 \\ 3.3 \end{bmatrix} \right) \left(\begin{bmatrix} 2 & 3 \end{bmatrix} - \begin{bmatrix} 2 & 3.3 \end{bmatrix} \right) \\ &+ \left(\begin{bmatrix} 3 \\ 4.9 \end{bmatrix} - \begin{bmatrix} 2 \\ 3.3 \end{bmatrix} \right) \left(\begin{bmatrix} 3 & 4.9 \end{bmatrix} - \begin{bmatrix} 2 & 3.3 \end{bmatrix} \right) \\ &= \begin{bmatrix} 1 & 1.3 \\ 1.3 & 1.69 \end{bmatrix} + \begin{bmatrix} 0 & 0 \\ 0 & 0.09 \end{bmatrix} + \begin{bmatrix} 1 & 1.6 \\ 1.6 & 2.56 \end{bmatrix} \\ &= \begin{bmatrix} 2 & 2.9 \\ 2.9 & 4.34 \end{bmatrix} \end{aligned}$$

Two-class linear discriminant analysis: example

Next,

$$S_2 = \begin{bmatrix} 2 & 2.9 \\ 2.9 & 4.34 \end{bmatrix}$$

and

$$S_w = \begin{bmatrix} 4 & 5.8 \\ 5.8 & 8.68 \end{bmatrix}$$

Thus, the optimal linear discriminant is

$$\mathbf{w} \propto S_w^{-1}(\mathbf{m}_1 - \mathbf{m}_2) = (-13.4074, 9.0741)^\top \propto (-0.8282, 0.5605)^\top$$

Two-class linear discriminant analysis: example

And the projections onto the linear discriminant for


Class 1 are

$$\{0.2928, 0.0252, 0.2619\}$$

and Class 2 are

$$\{-1.0958, -1.3635, -1.1267\}$$

.



USPS digits data

A two-dimensional PCA does not separate the digits 4 and 9 well:

PCA on digits 4 and 9



USPS digits data


```
library(IMIFA)
library(ggsci)
# Load USPS data
data("USPSdigits")
usps <- list()
usps$data <-
  (as.matrix(rbind(USPSdigits$train[, -1],
                   USPSdigits$test[, -1])) - -1) * 255 / 2
usps$label <- c(USPSdigits$train[, 1], USPSdigits$test[, 1])
usps_4_9 <- usps$data[usps$label %in% c(4, 9), ]

# Perform PCA
pca_4_9 <- prcomp(usps_4_9)
```

USPS digits data

```
plot_dat <-  
  data.frame(PC1 = pca_4_9$x[, 1],  
             PC2 = pca_4_9$x[, 2],  
             label = usps$label[usps$label %in% c(4, 9)])  
ggplot(plot_dat, aes(x = PC1, y = PC2, color = as.factor(label))) +  
  geom_point(alpha = .75, size = 2) +  
  theme_minimal() +  
  scale_color_d3(name = "digit", scale_name = "category10") +  
  ggtitle("PCA on digits 4 and 9") +  
  theme(aspect.ratio = 1,  
        panel.grid.major = element_blank(),  
        panel.grid.minor = element_blank(),  
        axis.text = element_blank(),  
        axis.title = element_blank())
```

USPS digits data



USPS digits data

```
library(MASS)
pca_95_idx <- which(cumsum(pca_4_9$sdev^2) / sum(pca_4_9$sdev^2) < .95)
pca_4_9_df <- as.data.frame(pca_4_9$x[, pca_95_idx]) |>
  mutate(label = usps$label[usps$label %in% c(4, 9)])
lda_4_9 <- lda(label ~., pca_4_9_df)

# compute projections
projections <-
  as.vector(pca_4_9$x[, pca_95_idx] %*% lda_4_9$scaling)
```

USPS digits data

```
plot_dat <- data.frame(  
  index = 1:nrow(usps_4_9),  
  a = projections,  
  label = usps$label[usps$label %in% c(4, 9)]  
)  
ggplot(plot_dat, aes(x = index, y = a, color = as.factor(label))) +  
  geom_point(alpha = .75, size = 2) +  
  theme_minimal() +  
  scale_color_d3(name = "digit", scale_name = "category10") +  
  ggtitle("LDA on digits 4 and 9 using PCs as features") +  
  ylab("Projection onto linear discriminant") +  
  theme(panel.grid.major = element_blank(),  
        panel.grid.minor = element_blank())
```

What if we have $c \geq 3$ classes? We can use the same idea, projecting data such that:

1. Points in the same class are tightly clustered together (minimize within-class variation)
2. The centers of the classes are as far apart from each other as possible (maximize between-class separation).

Multiple class linear discriminant analysis

We can still use the **total within-class scatter** matrix to describe the within-class variation:

$$\sum_{j=1}^c s_j^2 = \sum_j \mathbf{w}^\top S_j \mathbf{w} = \mathbf{w}^\top S_w \mathbf{w}$$

where

$$S_w = \sum_j S_j = \sum_j \sum_{i \in C_j} (\mathbf{x}_i - \mathbf{m}_j)(\mathbf{x}_i - \mathbf{m}_j)^\top$$

Multiple class linear discriminant analysis

For multiple class LDA, we define the **between-class scatter** as follows:

$$\sum_{j=1}^c n_j (\mu_j - \mu)^2 \quad \text{where} \quad \mu = \frac{1}{n} \sum_{j=1}^c n_j \mu_j \quad (\text{projected center})$$

Consider that

$$\mu = \frac{1}{n} \sum_{j=1}^c n_j (\mathbf{w}^\top \mathbf{m}_j) = \mathbf{w}^\top \left(\sum_{j=1}^c n_j \mathbf{m}_j \right) = \mathbf{w}^\top \left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \right) = \mathbf{w}^\top \mathbf{m}$$

where \mathbf{m} is the global centroid of $\mathbf{x}_1, \dots, \mathbf{x}_n$.

Multiple class linear discriminant analysis

The **between-class scatter** can be further simplified:

$$\begin{aligned}\sum_{j=1}^c n_j (\mu_j - \mu)^2 &= \sum_j n_j (\mathbf{w}^\top (\mathbf{m}_j - \mathbf{m}))^2 \\ &= \sum_j n_j (\mathbf{w}^\top (\mathbf{m}_j - \mathbf{m})) (\mathbf{m}_j - \mathbf{m})^\top \mathbf{w} \\ &= \mathbf{w}^\top \left(\sum_j n_j (\mathbf{m}_j - \mathbf{m}) (\mathbf{m}_j - \mathbf{m})^\top \right) \mathbf{w} \\ &= \mathbf{w} S_b \mathbf{w}\end{aligned}$$

where

$$S_b = \sum_j n_j (\mathbf{m}_j - \mathbf{m}) (\mathbf{m}_j - \mathbf{m})^\top$$

is the new between-class scatter matrix.

Multiple class linear discriminant analysis

This again yields a generalized Rayleigh quotient problem:

$$\max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{\sum_j n_j (\mu_j - \mu)^2}{\sum_j s_j^2} = \max_{\mathbf{w}: \|\mathbf{w}\|=1} \frac{\mathbf{w}^\top S_b \mathbf{w}}{\mathbf{w}^\top S_w \mathbf{w}}$$

Again, if S_w is nonsingular, then the solution the generalized eigenvector \mathbf{w}_1 of (S_b, S_w) :

$$S_b \mathbf{w} = \lambda_1 S_w \mathbf{w}_1 \iff S_w^{-1} S_b \mathbf{w}_1 = \lambda_1 \mathbf{w}_1$$

where λ_1 is the largest generalized eigenvalue of (S_b, S_w) .

Multiple class linear discriminant analysis

Note that now, we cannot simply use $\mathbf{w} \propto S_w^{-1}(\mathbf{m}_1 - \mathbf{m}_2)$:

$$\lambda_1 S_w \mathbf{w}_1 = S_w^{-1} S_b \mathbf{w}_1 = \sum_{j=1}^c n_j S_w^{-1} (\mathbf{m}_j - \mathbf{m})(\mathbf{m}_j - \mathbf{m})^\top \mathbf{w}_1$$

so all we know is that

$$\mathbf{w}_1 \in \text{Span}\{S_w^{-1}(\mathbf{m}_1 - \mathbf{m}), \dots, S_w^{-1}(\mathbf{m}_c - \mathbf{m})\}$$

Thus, \mathbf{w}_1 must be computed by solving the generalized eigenvalue problem.

Multiple class linear discriminant analysis

Recall that for the two-class problem, S_b was rank 1, so there was only one non zero eigenvalue. It turns out that

$$\text{rank}(S_w^{-1}S_b) = \text{rank}(S_b) \leq c - 1 \quad (\text{exercise})$$


so that LDA can be used to find at most a $c - 1$ -dimensional space upon which to project the data.

Multiple-class linear discriminant analysis: summary

1. Center and scale the data and find the class centroids $\mathbf{m}_1, \dots, \mathbf{m}_c$.
2. Compute S_w (within-class scatter) and S_b (between-class scatter).
3. Solve the generalized eigenvalue problem $S_b \mathbf{w} = \lambda S_w \mathbf{w}$ and find all non-zero eigenvalues.
4. Project the data onto the corresponding eigenvector(s).

USPS digits data: Multiple classes

PCA on digits 4, 7, and 9




USPS digits data: Multiple classes

```
usps_4_7_9 <- usps$data[usps$label %in% c(4, 7, 9), ]
pca_4_7_9 <- prcomp(usps_4_7_9)
plot_dat <-
  data.frame(PC1 = pca_4_7_9$x[, 1],
             PC2 = pca_4_7_9$x[, 2],
             label = usps$label[usps$label %in% c(4, 7, 9)])
ggplot(plot_dat,
       aes(x = PC1, y = PC2, color = as.factor(label))) +
  geom_point(alpha = .75, size = 2) +
  theme_minimal() +
  scale_color_d3(name = "digit", scale_name = "category10") +
  ggtitle("PCA on digits 4, 7, and 9") +
  theme(aspect.ratio = 1,
        panel.grid.major = element_blank(),
        panel.grid.minor = element_blank(),
        axis.text = element_blank(),
        axis.title = element_blank())
```

USPS digits data: Multiple classes

LDA on digits 4, 7, and 9 using PCs as features



USPS digits data: Multiple classes

```
pca_95_idx <-  
  which(cumsum(pca_4_7_9$sdev^2) /  
        sum(pca_4_7_9$sdev^2) < .95)  
pca_4_7_9_df <-  
  as.data.frame(pca_4_7_9$x[, pca_95_idx]) |>  
  mutate(label = usps$label[usps$label %in% c(4, 7, 9)])  
lda_4_7_9 <- lda(label ~., pca_4_7_9_df)  
ld1 <- as.vector(pca_4_7_9$x[, pca_95_idx] %*% lda_4_7_9$scaling[,1])  
ld2 <- as.vector(pca_4_7_9$x[, pca_95_idx] %*% lda_4_7_9$scaling[,2])
```

USPS digits data: Multiple classes

```
plot_dat <- data.frame(  
  index = 1:nrow(usps_4_7_9),  
  ld1 = ld1,  
  ld2 = ld2,  
  label = usps$label[usps$label %in% c(4, 7, 9)]  
)  
ggplot(plot_dat, aes(x = ld1, y = ld2, color = as.factor(label))) +  
  geom_point(alpha = .75, size = 2) +  
  theme_minimal() +  
  scale_color_d3(name = "digit", scale_name = "category10") +  
  ggtitle("LDA on digits 4, 7, and 9 using PCs as features") +  
  xlab("LD1") + ylab("LD2") +  
  theme(panel.grid.major = element_blank(),  
        panel.grid.minor = element_blank())
```

Darlingtonia data

```
darl <-  
  read.csv(  
    paste0("https://raw.githubusercontent.com/",  
           "jon-bakker/",  
           "appliedmultivariatestatistics/main/",  
           "Darlingtonia_GE_Table12.1.csv"),  
    header = TRUE  
  )
```

Darlingtonia data

```
options(width = 80)
head(darl)
```

	site	plant	height	mouth.diam	tube.diam	keel.diam	wing1.length	wing2.length
1	TJH	1	654	38.4	16.6	6.4	85	76
2	TJH	2	413	22.2	17.2	5.9	55	26
3	TJH	3	610	31.2	19.9	6.7	62	60
4	TJH	4	546	34.4	20.8	6.3	84	79
5	TJH	5	665	30.5	20.4	6.6	60	51
6	TJH	6	665	33.6	19.5	6.6	84	66

	wingsprea	hoodmass.g	tubemass.g	wingmass.g
1	55	1.38	3.54	0.29
2	60	0.49	1.48	0.06
3	78	0.60	2.20	0.16
4	95	1.12	2.95	0.24
5	30	0.67	3.36	0.08
6	82	1.27	4.05	0.21

Darlingtonia data

```
darl_data <- scale(darl[,3:ncol(darl)])  
darl_da <- lda(x = darl_data, grouping = darl$site)  
darl_da <- lda(darl$site ~ darl_data) # equivalent  
darl_da$prior
```


	DG	HD	LEH	TJH
	0.2873563	0.1379310	0.2873563	0.2873563

```
darl_da$scaling
```

	LD1	LD2	LD3
darl_dataheight	1.40966787	0.2250927	-0.03191844
darl_datamouth.diam	-0.76395010	0.6050286	0.45844178
darl_datatube.diam	0.82241013	0.1477133	0.43550979
darl_datakeel.diam	-0.17750124	-0.7506384	-0.35928102
darl_datawing1.length	0.34256319	1.3641048	-0.62743017
darl_datawing2.length	-0.05359159	-0.5310177	-1.25761674
darl_datawingsprea	0.38527171	0.2508244	1.06471559
darl_datahoodmass.g	-0.20249906	-1.4065062	0.40370294
darl_datatubemass.g	-1.58283705	0.1424601	-0.06520404
darl_datawingmass.g	0.01278684	0.0834041	0.25153893


Darlingtonia data

```
plot(darl_da, col = match(darl$site, unique(darl$site)))
```



Darlingtonia data

```
plot(darl_da, dimen = 1, type = "both")
```



Darlingtonia data

```
darl_predict <- predict(darl_da)
darl_table <- table(darl$site, darl_predict$class)
darl_table
```

	DG	HD	LEH	TJH
DG	18	0	2	5
HD	0	11	1	0
LEH	3	0	21	1
TJH	2	0	3	20

How might dimension reduction methods be helpful in prediction?

Examples:

- Matrix completion (Netflix prize)
- Image denoising
- Timeseries forecasting
- Spatial interpolation (kriging)

In all of these cases, we predict a large data matrix/vector whose entries are only partially or noisily observed.

